

HPC 565: Financial Mathematics - Asian Option Pricing

Dominik Dahlem

May 9, 2008

1 Introduction

Asian Options rely on the average of the underlying asset over a predetermined averaging period leading up to the maturity time T . Those options are used when price stability of the underlying asset is particularly important.

This report outlines the approach and the results of Asian option pricing. The evolution of the share price follows a lognormal model given as

$$S_n = S_{n-1} e^{(r - \frac{1}{2}\sigma^2 + \sigma N_n(0,1))} \quad (1)$$

where $N_n(0, 1), n = 1, 2, \dots$ is a sequence of independent standard normal random variables. r is specified as the daily risk-free interest rate and σ defines the daily volatility. If the annual volatility is given, then the volatility has to be transformed given Equation 2:

$$\sigma_{daily} = \sigma_{annum} \sqrt{1/252} \quad (2)$$

where 252 is the number of business days per year.

Given the daily share prices in the interval $[t, T]$ the Asian call option maturing at time T is given by the payoff function

$$\max(0, S_{ave} - K) \quad (3)$$

where S_{ave} is the average value of the underlying share price over a specified period of time. Here, the time period taken into account cover the last 21 days before maturity of the option:

$$S_{ave} = \frac{1}{22} \sum_{l=0}^{21} S_{T-l} \quad (4)$$

The price $C(t)$ of an Asian Call option maturing at time T is then

$$C(t) = e^{-r(T-t)} \mathbb{E}[(S_{ave} - K)^+] \quad (5)$$

where K is the strike price, $e^{-r(T-t)}$ is the discount factor, and

$$X^+ = \begin{cases} X, & \text{if } X \geq 0 \\ 0, & \text{otherwise} \end{cases} \quad (6)$$

The price $C(t)$ can be found using N independent price evolutions $S_{n=T-21}^{(i)T}$, $i = 1, 2, \dots, N$. The Asian Call option price estimator is then given by

$$\hat{C}(t) = e^{-r(T-t)} \frac{1}{N} \sum_{i=1}^N [(S_{ave}^{(i)} - K)^+] \quad (7)$$

2 Implementation Details

The simulation is implemented using R¹ as an independent R package. Each R routine is documented with the R documentation environment. Once the library is loaded into an R session help pages for the individual methods can be displayed using `?method.name`. Or, the package documentation can be displayed using `package?asianOptionPricing`. This documentation pages provides an overview of the package and the implemented routines. Further, it gives some examples of how to use this package. These examples are taken from simulation script which is supplied with the project distribution.

The layout of the project is as follows:

- src (the sources)
 - R (the main R script)
 - package (the R package)
- doc (report)
 - images (the graphs)

2.1 Parallel Execution

The application was implemented with parallel execution in mind. The distribution can be configured (see Section 2.2 for more details on that) for parallel and serial use.

¹<http://www.r-project.org>

R provides some efficient primitives to apply a function to a sequence of values. These functions were used to run the Monte Carlo simulation with a number of different configuration parameters, i.e., over a specified range of the annual volatility or over a specified range of share prices to simulate the Delta and Gamma values. Usually, these are implemented in for-loops, e.g., in C/C++. These primitives, such as **apply** and **lapply**, lend themselves to being executed in parallel, because they do not have overlapping contexts. The R library snow provides a nice way of parallelising those primitives with or without load-balancing.

The R file `asian.R` implements the Asian call option pricing routines, which are part of the `asianOptionPricing` package. `price.option.sigma` and `price.option.S` are the two methods that call the underlying Monte Carlo simulation with the **lapply** primitive. If a valid cluster parameter is passed into those methods the `clusterApplyLB` routine from the snow package is used to parallelise the execution of the Monte Carlo simulation. This method is load-balanced, which means that the sequence of values for which the simulation is run, can be larger than nodes being available in the cluster. That means, snow will take care of distributing the load to the available nodes. Otherwise, if the cluster parameter is not specified, the serial version is called instead.

The `asianOptionPricing` package itself does not set up the cluster. Instead, the supplied R script – the driver of the simulation – is expected to set up a clustered environment and pass the cluster configuration into the simulation methods mentioned above. Consequently, the option pricing package can be installed and used for both parallel and serial execution transparently.

Another advantage of snow is that it hides a lot of the low-level details of parallel implementations. It supports the Message Passing Interface, Parallel Virtual Machines, and sockets in a transparent way. The simulation script `asian_pricing.R` sets up an MPI cluster, which requires the PBS script to prepare the parallel execution of this script. These steps include generating a configuration file with the nodes listed which participate in the cluster and starting `mpd`. Once `mpd` is running the `mpiexec` is called with the parameter `"-n 1"` which spawns R and advises R to take control of `mpi_comm_spawn`. When the simulation is finished `mpd` is stopped in the PBS script.

The PBS script creates a text file called `NODEFILE` which contains all participating nodes in the cluster setup. The R simulation script reads this file and calls `makeCluster(numNodes, type="MPI")`, where `numNodes` is the number of nodes as specified in the `NODEFILE`. This avoids the need of hard-coding the MPI cluster setup in the R script.

As a measure of controlling exceptions in the execution flow of the R script `asian_pricing.R`, the `.Last` routine was implemented to catch any such runtime errors and clean up the cluster before exiting from the script. This method will

always be called before executing the simulation script, whether an exception occurred or not.

2.2 Configuration and Installation

The support for a parallel execution is provided via a configure script. If it is enabled, configure will generate the R script file with the respective snow API calls enabled. Otherwise, they are disabled. Also, the DESCRIPTION file as part of the Asian option pricing package is generated to include the snow library as a dependency. When the Asian option pricing library is installed, these dependencies have to be available in the local R installation.

Additionally, the report can be generated using the build system.

The assignment ships with a configure script generated by autotools. The installation procedure requires the following steps which are outlined in more detail below. First configure the package (i.e., for serial or parallel use), install the Asian option pricing package, run the simulation script.

2.2.1 Configure

The following configure options are supported:

- **–enable-mpi**: enables the snow library for parallel execution.
- **–enable-report**: enables the report generation.

So, to enable parallel execution change into the project root directory and do:

```
> ./configure --enable-mpi
```

Otherwise, call `./configure` without any arguments for serial use.

To generate the report PDF file, enable the report generation feature with `configure` and call `make` in the root directory of the project distribution. In this case `configure` will check whether `latex` and the respective tools are installed.

2.2.2 Installation

Apart from a basic R (≥ 2.6) installation, this package requires the `fOptions` and `gplots` packages. If those packages are not installed yet open an R session and install them using the R interface

```
> install.packages("fOptions", lib=Sys.getenv("R_LIBS_USER"), \
  dependencies=TRUE, method="wget")
> install.packages("gplots", lib=Sys.getenv("R_LIBS_USER"), \
  dependencies=TRUE, method="wget")
```

The method parameter specifies the usage of the wget utility to download the package. The advantage of this approach is that it uses the system environment variables `http_proxy` if configured to go through the network proxy.

`fOptions` provides methods to generate normal random numbers using an underlying Sobol sequence or a pseudo random number generator. The Asian option pricing package provides a configuration option to select either one of those.

In order to install the Asian option pricing package as a user library you would first need to find out which user library path is configured as the default one with:

1. `> R`
2. `> Sys.getenv("R_LIBS_USER")` (in the R session)

On my computer the default path for user libraries is `" /R/i686-pc-linux-gnu-library/2.7"`. Create these directories, if they don't exist and then install the library only for the user with

1. `> cd src/package`
2. `R CMD INSTALL -l ~/R/i686-pc-linux-gnu-library/2.7 \`
`asianOptionPricing`

Alternatively, the dependencies and the option pricing package can be installed system-wide:

```
> R (start an R session)
> install.packages("fOptions", dependencies=TRUE, method="wget")
> install.packages("gplots", dependencies=TRUE, method="wget")
> q() (quit the R session)
> cd src/package
> R CMD INSTALL asianOptionPricing
```

2.2.3 Execution

Now, once the Asian option pricing library is installed, the simulation script can be executed in the root directory of the project distribution with:

```
> R -f src/R/asian_pricing.R
```

Or, if this project was configured for parallel execution with `--enable-mpi`, the job can be submitted to a cluster with `qsub option_pricing.pbs`. The number of nodes can be adjusted in the PBS script (default is 4 nodes).

Each execution will create a timestamped directory in the `./eval` folder of the root project directory. This is to ensure that many simulation runs will not overwrite their generated data. The timestamped directory is per R session. Therefore, it is best to call R from the command-line.

3 MC Pricing of Asian Options

Given the Equation 7, a Monte Carlo routine can be employed to estimate the payoff and consequently the discounted Asian Call option price.

The MC algorithm implements the following steps:

1. Evolve $S_{n=T-21}^{(i)T}$
2. Calculate the average share price S_{ave} in the interval $[T - 21, T]$
3. Calculate the payoff $\max(0, S_{ave} - K)$
4. Repeat $\forall i = 1, 2, \dots, N$
5. Calculate $\mathbb{E}(S_{ave} - K)^+ \approx \frac{1}{N} \sum_{i=1}^N \left[(S_{ave}^{(i)} - K)^+ \right]$

Figure 1 presents a sample path generated using 1000 iterations.

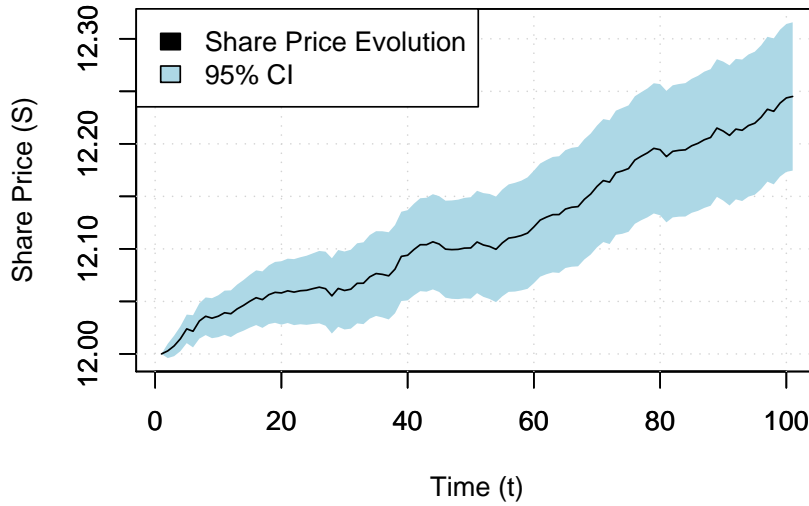


Figure 1: Share Price Evolution

The Monte Carlo routine reports the standard deviation which can be used to plot the confidence intervals for derived values, such as the discounted option price at $t = 0$, and the respective Δ and Γ sensitivities. For all results provided in Section 4 the 95% confidence intervals are reported using the standard error defined as $s.e. = \frac{s.d.}{\sqrt{N}}$, where N is the number of MC steps.

The Monte Carlo routine exposes two configuration options which determine the approach to the random number generation. The first one `pseudo` is a logical value which switches between a pseudo random number generator or a Sobol sequence as the underlying random number stream for normal variates. The second one `init` is also a logical value that determines whether the random number generator should be initialised before the MC steps. For the purpose of the simulations described in this report, `init` was set to `TRUE`, because the Monte Carlo routine is run multiple times for different option pricing parameters. Consequently each MC run always generates the same stream.

The Figure 2 shows the price evolution using an underlying Sobol sequence of the lognormal model. The price evolution is a bit more choppy, because only 100 iterations were used to calculate the mean path. Also, this graph shows that the confidence interval is narrower compared to Figure 1 with only 100 MC iterations. That is an order of an magnitude less than the simulation with pseudo random numbers.

A Sobol sequence is a low discrepancy sequence with space-filling properties. Compared to a pseudo random number this leads to a smaller error in the Monte Carlo routine. The construction of a low discrepancy sequence implies that an empirical standard deviation of the estimator in the Monte Carlo routine cannot be calculated, unless those sequences are themselves randomised. Here Sobol sequences are drawn in N dimensions. One dimension for each Monte Carlo iteration. Each dimension is scrambled to ensure that some of the determinism is removed. The seed for the random scrambling is the same as for the pseudo random number generation and so the sequences are retained between different Monte Carlo runs with different configurations.

4 Results

This section presents the results for five experiments conducted using the Monte Carlo Asian Call option pricing routine.

4.1 Volatility Bracket

Figure 3 and Figure 4 presents two graphs plotting the price of the Asian Call option against σ_{annum} . For this experiment the strike price is set to $K = 14$, time time-to-maturity $T = 100$, the daily risk-free interest rate $r = 0.0002$, and the share price at $t = 0$ is $S_0 = 12$. Plugging σ_{annum} into the lognormal model given in Equation 1 requires a transformation into the daily volatility using Equation 2.

It can be seen that the annual volatility bracket $0 \leq \sigma_{annum} \leq 600\%$ presents an increase in the Asian Call option price leading up to $\sigma_{annum} \approx 5.2$ after which

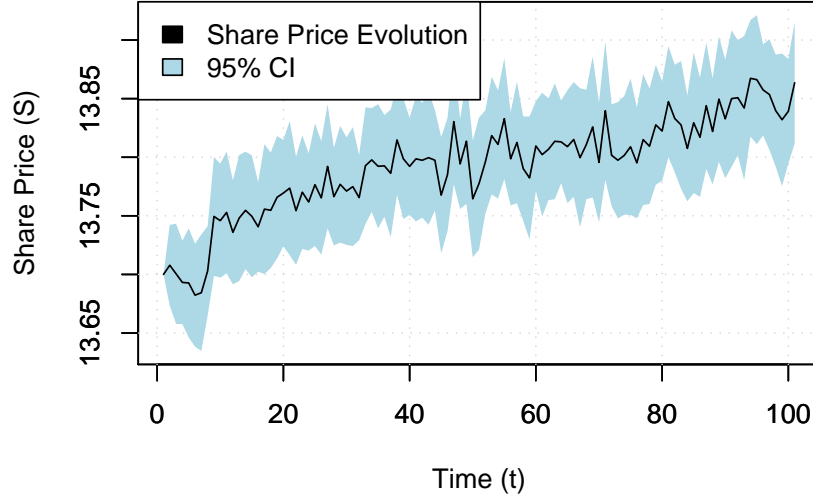


Figure 2: Share Price Evolution

the option price declines again. The maximum discounted option price at is $C(t = 0) \approx 9.3$.

4.2 Option Price

Figure 5 shows the plot of the Asian Call option price against the share price on the interval $0 \leq S_0 \leq K$.


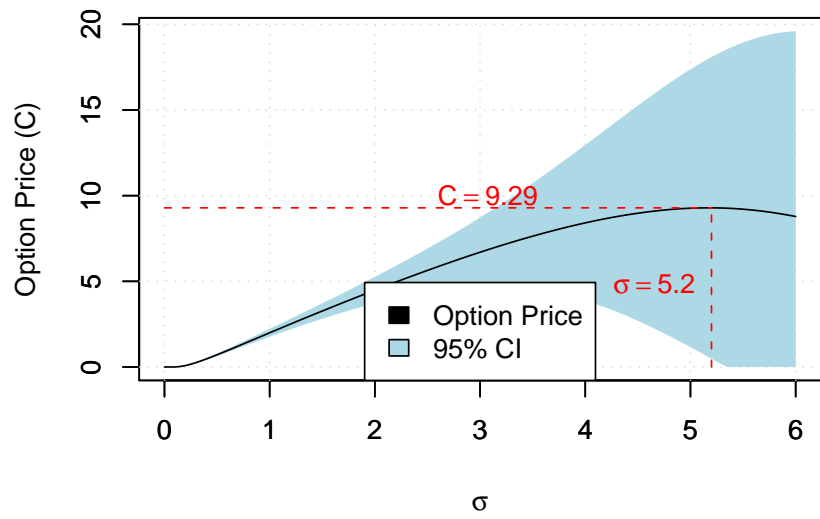
As expected, the option price is 0 leading up to $S_0 \approx 10$ after which it slowly increases to $S_0 \approx 12$ and then increasing linearly in S_0 . The rate of change is better captured in the Δ sensitivity plot Figure 6 in the next section.

4.3 Option Price Sensitivities

This section shows the plots of the Asian Call option sensitivities Δ and Γ .

Δ is defined as the rate of change of the option price with respect to the underlying share price (or more generally the underlying asset).

Approximately in the interval of $10 \leq S_0 \leq 18$ the rate of change is sub-linear before it reaches a value of $\Delta = 1$ (see Figure 6). The European Call option sensitivity Δ closely approximates the one for the Asian Call option price.

Figure 3: Asian Call option against $0 \leq \sigma_{\text{annum}} \leq 100\%$ Figure 4: Asian Call option against $0 \leq \sigma_{\text{annum}} \leq 600\%$


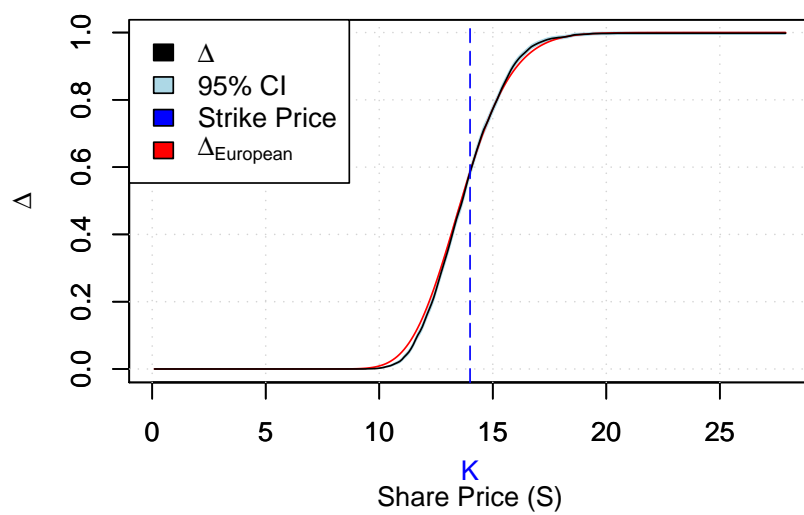


Figure 5: Asian Call option against the Share price

Figure 6: Asian Call option Δ against the Share price at $t = 0$

The gamma sensitivity shown in Figure 7 is a little choppy, because the variability induced through the random numbers is most pronounced in the second partial derivative of the option price with respect to the share price. Here, the analytical form of the European call option gamma approximates the Asian Call option one, too.

Hedging techniques purely based on the delta sensitivity introduces uncertainty because it measures the rate of change of the option price to the share price (i.e., the slope), while gamma measures the curvature of the relationship of the option price and the stock price. So when gamma is small the call/share price relationship tends to be linear, whereas a large gamma value implies a more pronounced curvature of the relationship. Consequently, the larger the gamma value is, the higher the hedging frequency ought to be in order to maintain a delta-neutral position.

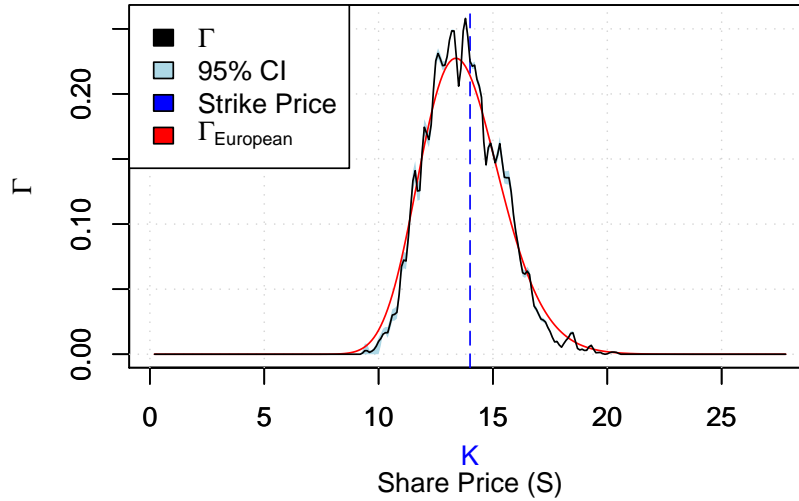


Figure 7: Asian Call option Γ against the Share price at $t = 0$

The delta and gamma values were calculated using the central-difference scheme, i.e.,

$$f'(x) = \frac{f(x+h) - f(x-h)}{2h} \quad (8)$$

Since no boundary conditions are given, each derivative requires $f(x+h)$ and $f(x-h)$ to be defined via Monte Carlo simulation. For the sensitivity simulation

the step size of the share price was set to $h = 0.1$, i.e., an increase of 10 cent per Monte Carlo simulation.

Since 2000 Monte Carlo iterations for each data point on the graphs was performed, the 95% confidence interval in the sensitivity plots is very narrow.

5 Simulation Results for Sobol Sequences

This section just shows the results obtained using scrambled Sobol sequences as the underlying stream for Normal random number generation. The curvature in the option price to share price plot is a lot less pronounced and therefore the Delta and Gamma are curves are narrower. I'm not sure why this is the case. The configuration used for this setup was

```
S0 <- 13.7
K <- 14
S.min <- 0
S.max <- 2 * K
S.step <- 0.1
r <- 0.0002
sigma <- 0.013
sigma.min <- 0
sigma.step <- 0.05
T <- 100
l <- 22
reps <- 100
seed <- 123574903
pseudoRnd <- FALSE
euro <- TRUE
```

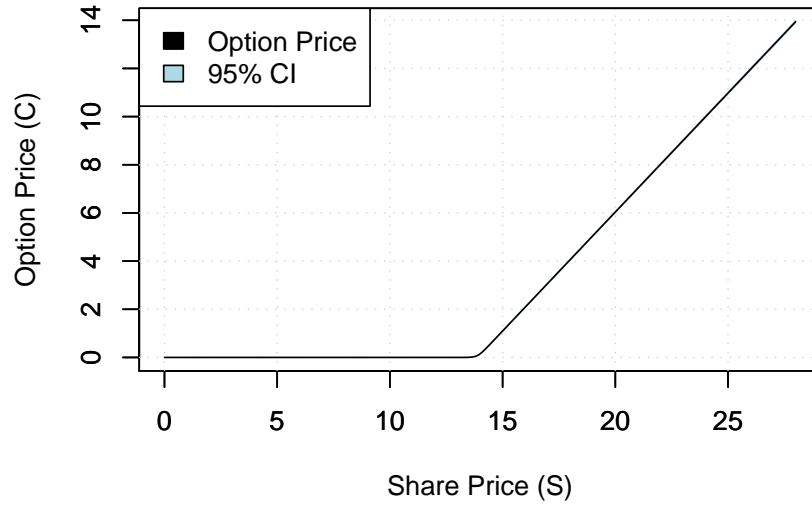


Figure 8: Asian Call option against the Share price using Sobol Sequences




Figure 9: Asian Call option Δ against the Share price at $t = 0$ using Sobol sequences

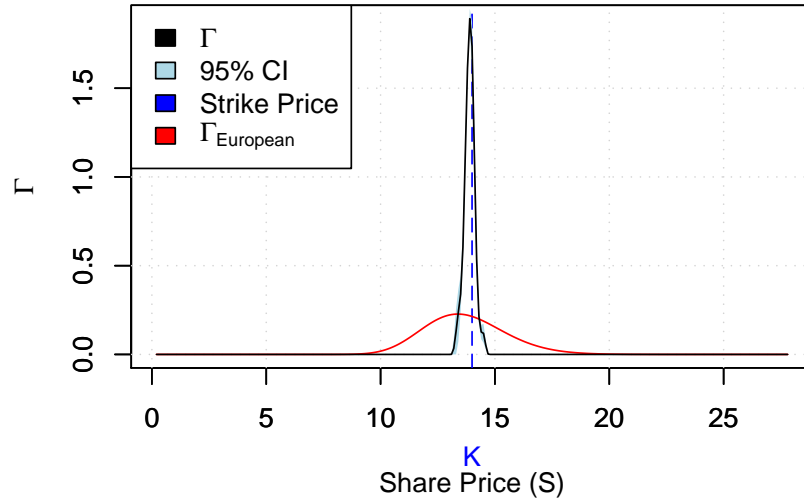


Figure 10: Asian Call option Γ against the Share price at $t = 0$ using Sobol sequences




Figure 11: Asian Call option against $0 \leq \sigma_{\text{annum}} \leq 100\%$ using Sobol sequences